RECENT GPT MODELS

A text-only reference guide for testing a Retrieval-Augmented Generation (RAG) project

Purpose: This document summarizes selected recent OpenAI GPT model announcements and their broad intended uses. It is written as a small knowledge base for PDF ingestion, chunking, retrieval, and question-answering experiments.

Information date: 13 October 2026

Note: Model names, availability, capabilities, and API details can change. This is a concise educational summary, not a complete model catalog or a substitute for current official documentation.

Contents

• 1. GPT-6 family overview

• 2. GPT-6.1 Sol

• 3. GPT-6 Astra

• 4. GPT-6 Sol and GPT-6 Luna

• 5. GPT-5.6 family context

• 6. How to choose a model

• 7. RAG testing questions

• 8. Official references

1. GPT-6 Family Overview

OpenAI announced GPT-6 in October 2026 and described a rollout in ChatGPT with an Intelligent UI experience. The GPT-6 family includes models designed for different balances of capability, cost, and workload complexity. The exact model available to a user can depend on the product, plan, region, and rollout status.

The current OpenAI API model catalog describes GPT-6 Astra as the flagship option for demanding reasoning and coding, GPT-6.1 Sol as a lower-cost option aimed at complex work, and GPT-6 Luna as an efficient choice for high-volume tasks. Always check the current official model catalog before selecting a model for an application.

2. GPT-6.1 Sol

OpenAI introduced GPT-6.1 Sol on 29 September 2026. The announcement presents it as a model for coding, computer use, and professional work, with a lower token price than the Astra model. It is intended to offer strong capability for complex tasks where cost matters as well as quality.

Potential use cases

• Software development assistance and code review.

• Multi-step office and professional workflows.

• Tasks involving the use of tools or computer interfaces.

• Applications that need strong capability while managing inference costs.

3. GPT-6 Astra

OpenAI introduced GPT-6 Astra in September 2026 and described it as its most capable model for demanding work. The announcement highlights computer use, browsing, software engineering, cybersecurity, science, and professional tasks. A model's published capabilities do not guarantee that every answer will be correct; applications should still validate outputs and apply appropriate safety and quality checks.

When it may be appropriate

A flagship model may be appropriate when a task is difficult, mistakes are costly, or strong reasoning is more important than minimizing inference cost. For routine classification, extraction, or high-volume tasks, a smaller or less expensive model may be more practical.

4. GPT-6 Sol and GPT-6 Luna

OpenAI announced GPT-6 Sol and GPT-6 Luna on 22 September 2026. The API catalog positions Sol as a model for complex work with a balance of capability and cost, while Luna is positioned for cost-sensitive, high-volume workloads. These descriptions are broad selection guidance; a developer should evaluate candidate models on their own data and tasks.

In practice, a cost-efficient model may be suitable for short summaries, simple question answering, extraction, or routing. A stronger reasoning model may be worth the additional cost for complex analysis or difficult multi-step tasks. Model choice should be based on measured quality, latency, reliability, and total cost.

5. GPT-5.6 Family Context

OpenAI previewed GPT-5.6 Sol in June 2026, describing improvements aimed at coding, science, and cybersecurity, along with additional safeguards for higher-risk use. Model families and release stages can overlap, so a preview announcement does not necessarily mean that a model is generally available to every user.

For RAG developers, the key lesson is that a model's name alone does not determine application quality. The quality of the source documents, chunking strategy, retrieval results, prompt, and evaluation process also matter.

6. How to Choose a Model

Need	Starting point	What to check
Complex reasoning or coding	Flagship model such as GPT-6 Astra	Accuracy on difficult tasks, latency, cost
Strong capability with cost awareness	GPT-6.1 Sol or GPT-6 Sol	Quality versus price on your workload
High-volume, routine tasks	GPT-6 Luna	Throughput, consistency, and sufficient quality
Document Q&A with RAG	Test a suitable available model with retrieved context	Groundedness, citation correctness, retrieval quality

Important: Availability and API model IDs can change. Confirm the exact model identifier, capabilities, and pricing in official documentation before writing production code. API access may require a paid account even when ChatGPT itself offers a free plan.

7. Questions to Test Your RAG Chatbot

- Which GPT-6 model is positioned for the most demanding reasoning and coding work?
- What types of tasks are highlighted for GPT-6.1 Sol?
- How is GPT-6 Luna positioned in the API model catalog?
- When was GPT-6 Astra announced?
- What kinds of work were highlighted in the GPT-5.6 Sol preview?
- Why should a developer evaluate a model on their own task instead of choosing by name alone?
- What factors besides the LLM influence the quality of a RAG application?
- Does a model announcement guarantee that every user has access to that model?
- What should be checked before using a model in production?
- According to this document, which model would be a possible starting point for high-volume routine tasks?

8. Official References

OpenAI Research and Release Index
https://openai.com/research/index/release/

OpenAI API Models Catalog
https://developers.openai.com/api/docs/models

GPT-6 Astra announcement
https://openai.com/index/gpt-6-astra/

OpenAI Model Release Notes
https://help.openai.com/en/articles/9624314-model-release-notes

GPT-5.6 Sol preview announcement
https://openai.com/index/previewing-gpt-5-6-sol/

Prepared as a text-focused sample knowledge document for AI Document Assistant and RAG experimentation. The document contains selectable text rather than images, making it suitable for PDF text extraction and chunking tests.